

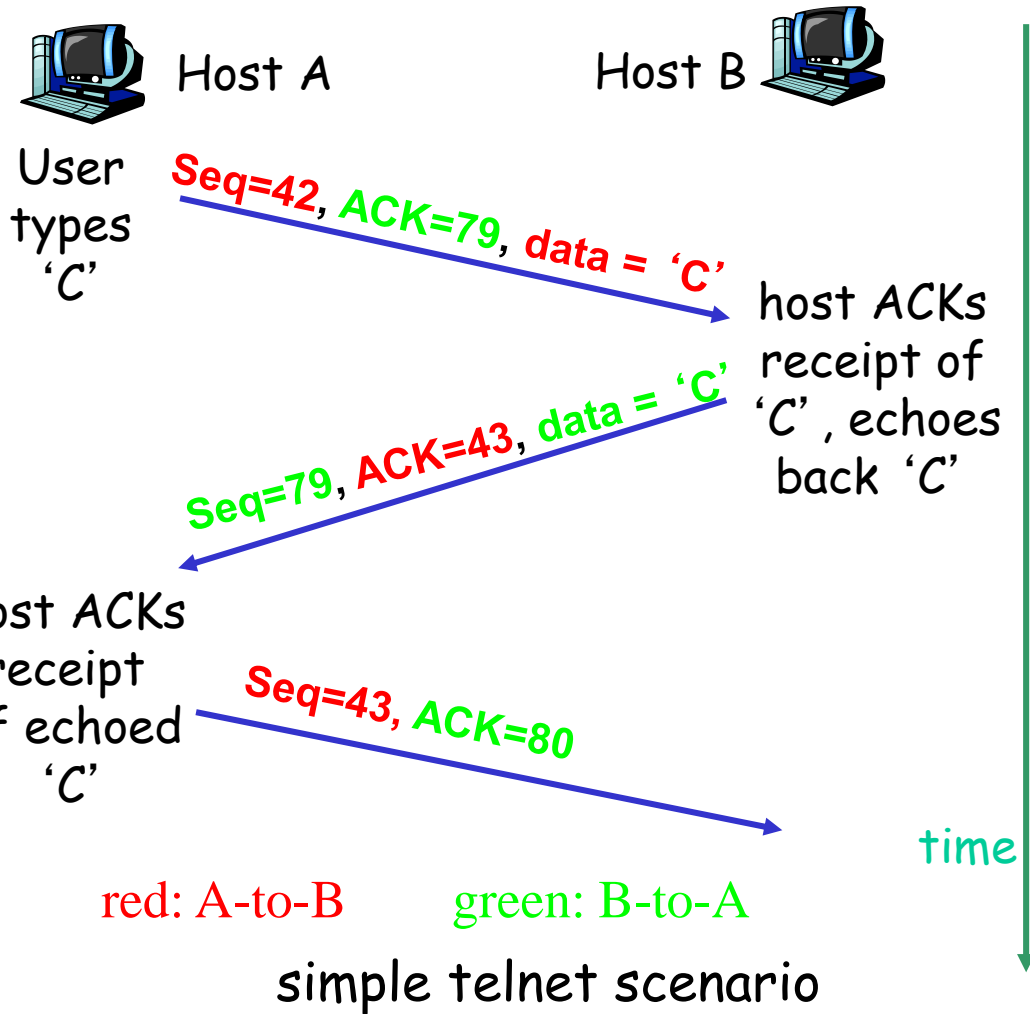
TCP Seq. #'s and ACKs

Seq. #'s:

byte stream
“number” of first
byte in segment's
data

ACKs:

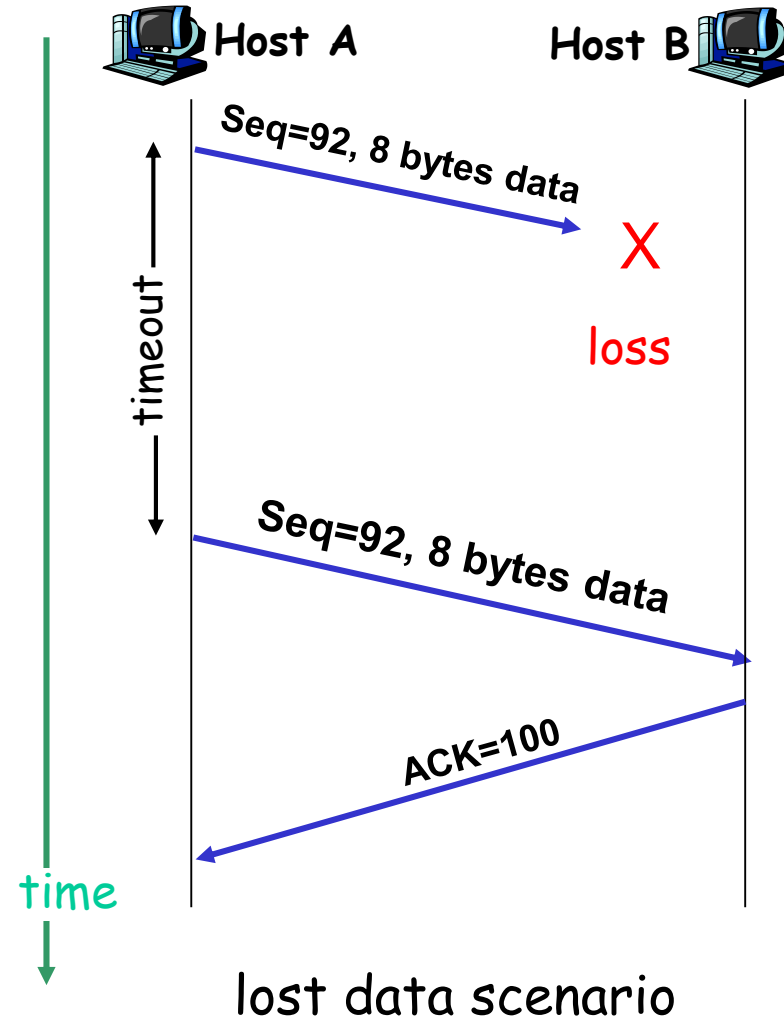
seq # of next byte
expected from
other side



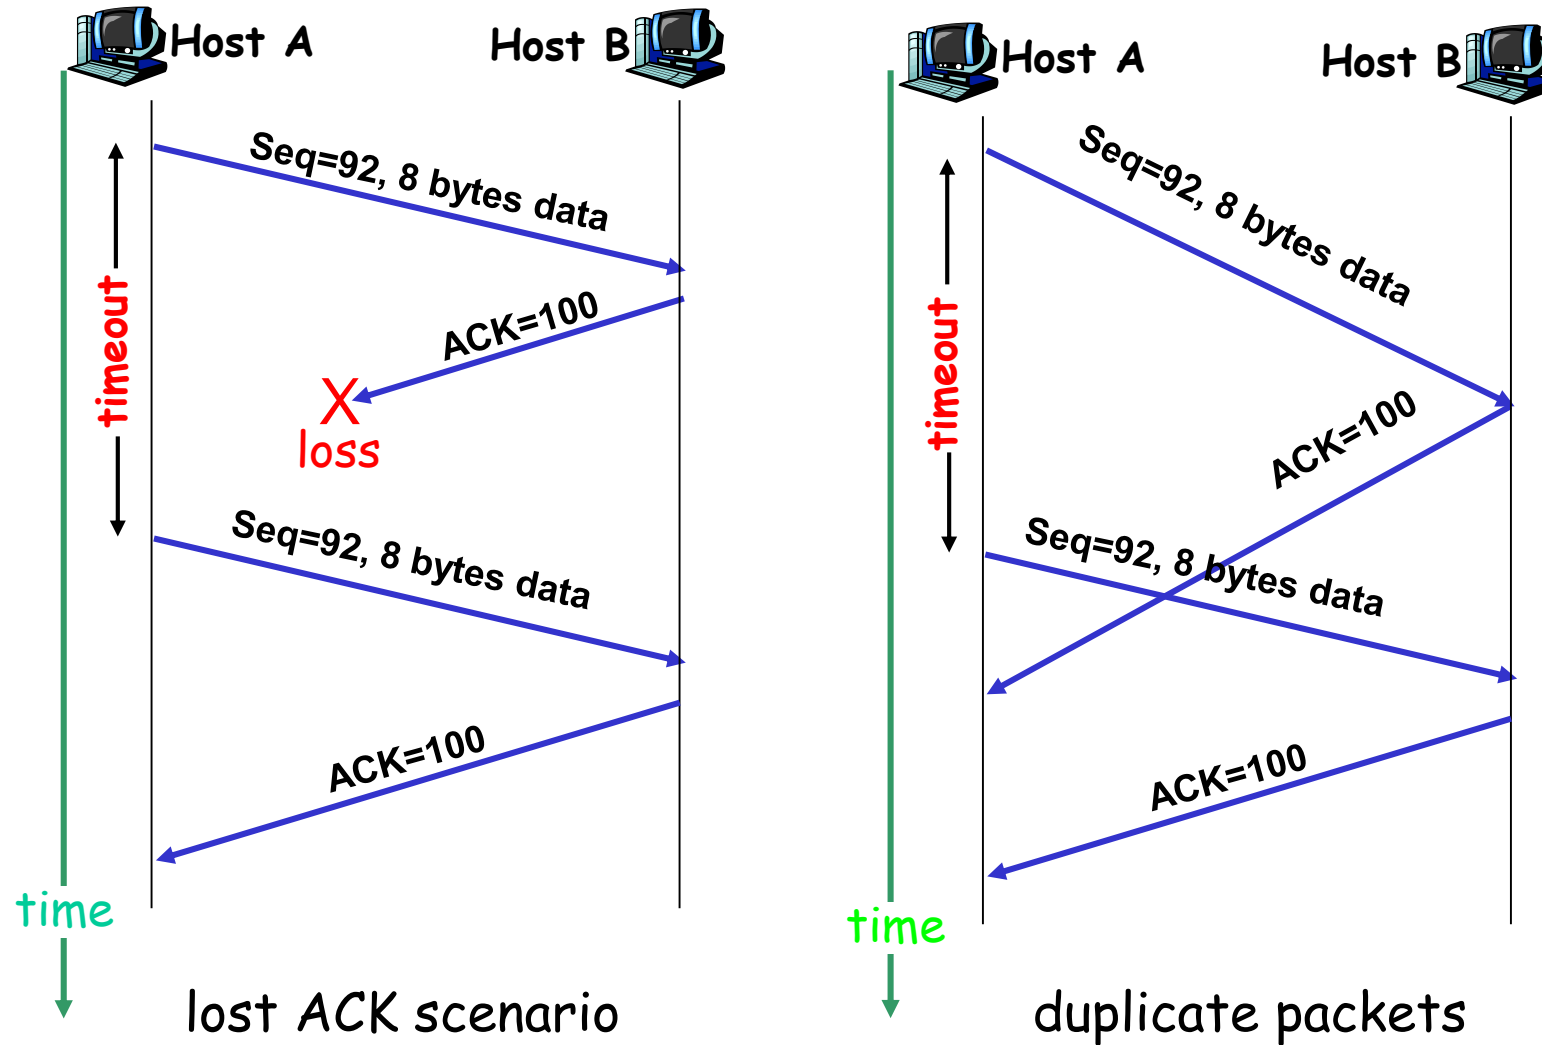
TCP: Error Scenarios

Questions for you:

- How to detect lost packets?
- How to “recover” lost packets?
- Potential consequence of retransmission?
- How to detect duplicate packets?
- “State” maintained at sender & receiver?



TCP: Error Scenarios (cont' d)



A Simple Reliable Data Transfer Protocol

“Stop & Wait” Protocol (aka “Alternate Bit” Protocol)

Sender algorithm:

- **Send Phase:** send data segment (n bytes) w/ **seq=x**, buffer data segment, set timer
- **Wait Phase:** wait for ack from receiver w/ **ack= x+n**
 - if received ack w/ **ack=x+n**, set **x:=x+n**, and go to sending phase with next data segment
 - if time out, resend data segment w/ **seq=x**.
 - if received ack w/ **ack != x+n**, ignore (or resend data segment w/ **seq=x**)

Receiver algorithm:

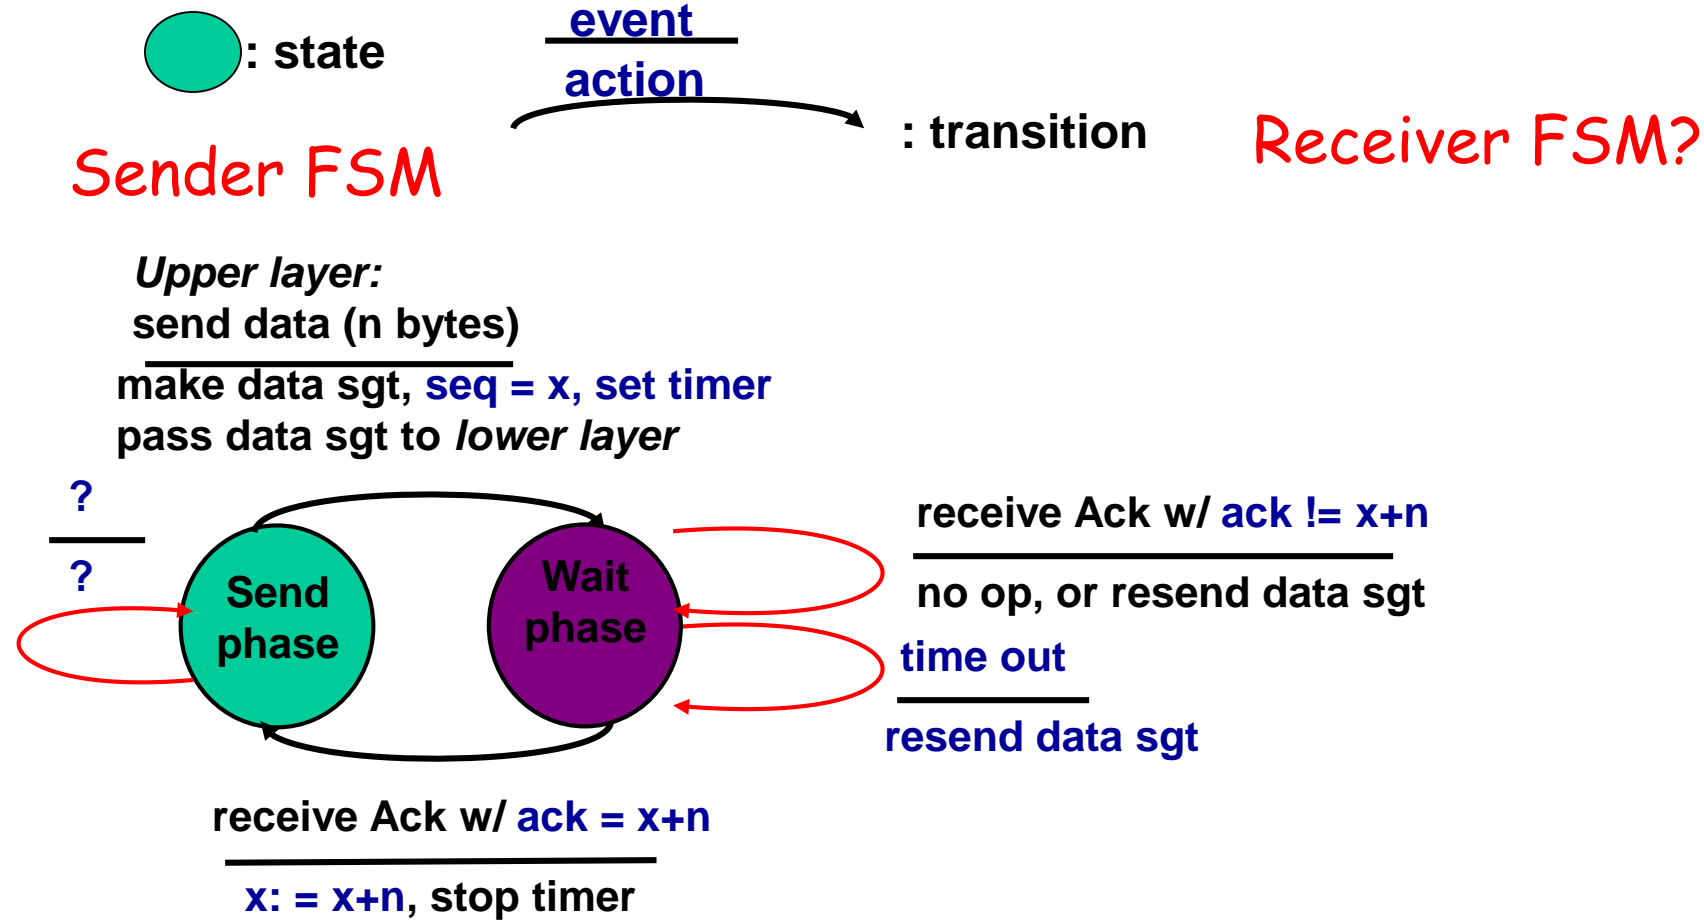
Wait-for-Data:

wait for data packet with the (expected) **next-seq = x**

- if received Data packet w/ **seq. =x** and of size n bytes: send ACK pkt w/ **ack = x+n**; **set next-seq:= x+n**; go back to “Wait-for-Data”;
- If received Data packet w/ **seq != x**, (re-)send ACK pkt w/ **ack= next-seq**; go back to “Wait-for-Data”;

Q: what is the “state” information maintained at the sender & receiver, resp.?

SRDTP: Finite State Machine



info ("state") maintained at sender:
phase it is in (*send*, or *wait*), ack expected, data sgt sent (seq #), timer

TCP Connection Set Up

Three way handshake:

TCP sender, receiver establish “connection” before exchanging data segments

- initialize TCP variables:
 - seq. #s
 - buffers, flow control info
- *client*: end host that initiates connection
- *server*: end host contacted by client

Step 1: client sends TCP **SYN** control segment to server


- specifies initial seq #

Step 2: server receives SYN, replies with **SYN+ACK** control segment

- ACKs received SYN
- specifies server → receiver initial seq. #

Step 3: client receives **SYN+ACK**, replies with **ACK** segment (which may contain 1st data segment)

TCP 3-Way Hand-Shake



Question:

- What kind of "state" client and server need to maintain?
- What initial sequence # should client (and server) use?

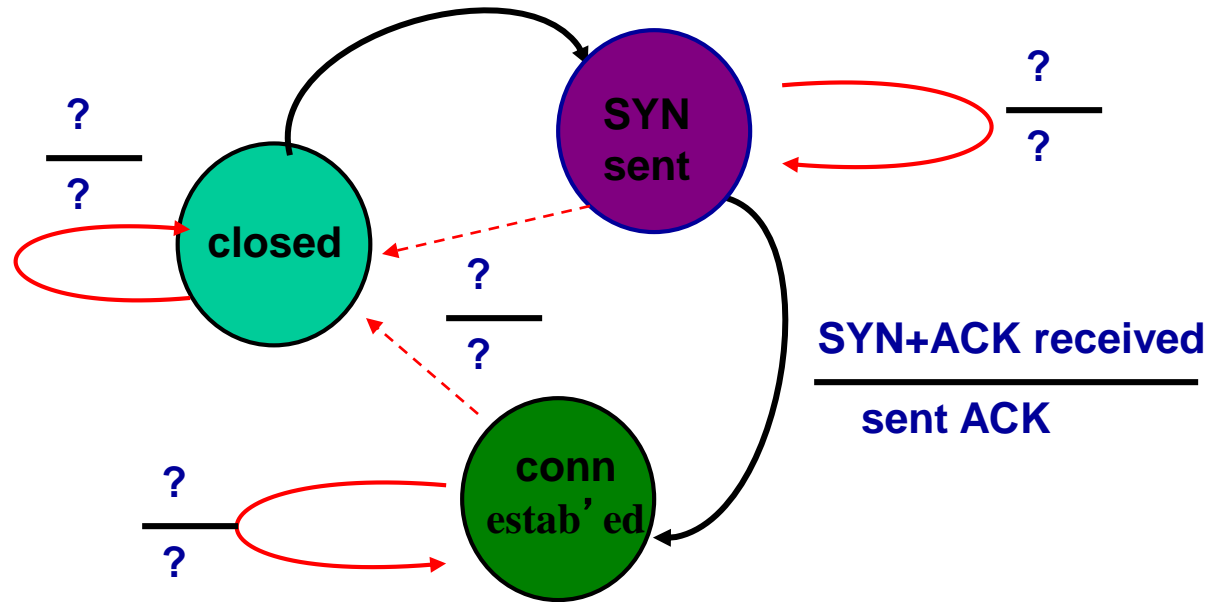
3-Way Handshake: Finite State Machine

Client FSM?

Server FSM?

Upper layer:
initiate connection

sent SYN w/ initial seq = x



info ("state") maintained at client?

Connection Setup Error Scenarios

- Lost (control) packets
 - What happen if SYN lost? client vs. server actions
 - What happen if SYN+ACK lost? client vs. server actions
 - What happen if ACK lost? client vs. server actions
- Duplicate (control) packets
 - What does server do if duplicate SYN received?
 - What does client do if duplicate SYN+ACK received?
 - What does server do if duplicate ACK received?

Connection Setup Error Scenarios (cont' d)

- Importance of **(unique)** initial seq. no.?
 - When receiving SYN, how does server know it's a new connection request?
 - When receiving SYN+ACK, how does client know it's a legitimate, i.e., a response to its SYN request?
- Dealing with **old duplicate** (aka “ghost”) packets from old connections (or from malicious users)
 - If not careful: “TCP Hijacking”
- How to choose unique initial seq. no.?
 - randomly choose a number (and add to last syn# used)
- Other security concern:
 - “SYN Flood” -- denial-of-service attack

TCP: Closing Connection

Remember TCP duplex connection!

Client wants to close connection:

Step 1: client end system sends TCP **FIN** control segment to server


Step 2: server receives FIN, replies with **ACK**. half closed

Step 3: client receives FIN.

half closed, wait for server to close

Server finishes sending data,
also ready to close:

Step 4: server sends FIN.



TCP: Closing Connection (revised)


Two Army Problem!

Step 5: client receives FIN, replies with ACK.

- Enters “timed wait” - will respond with ACK to received FINs

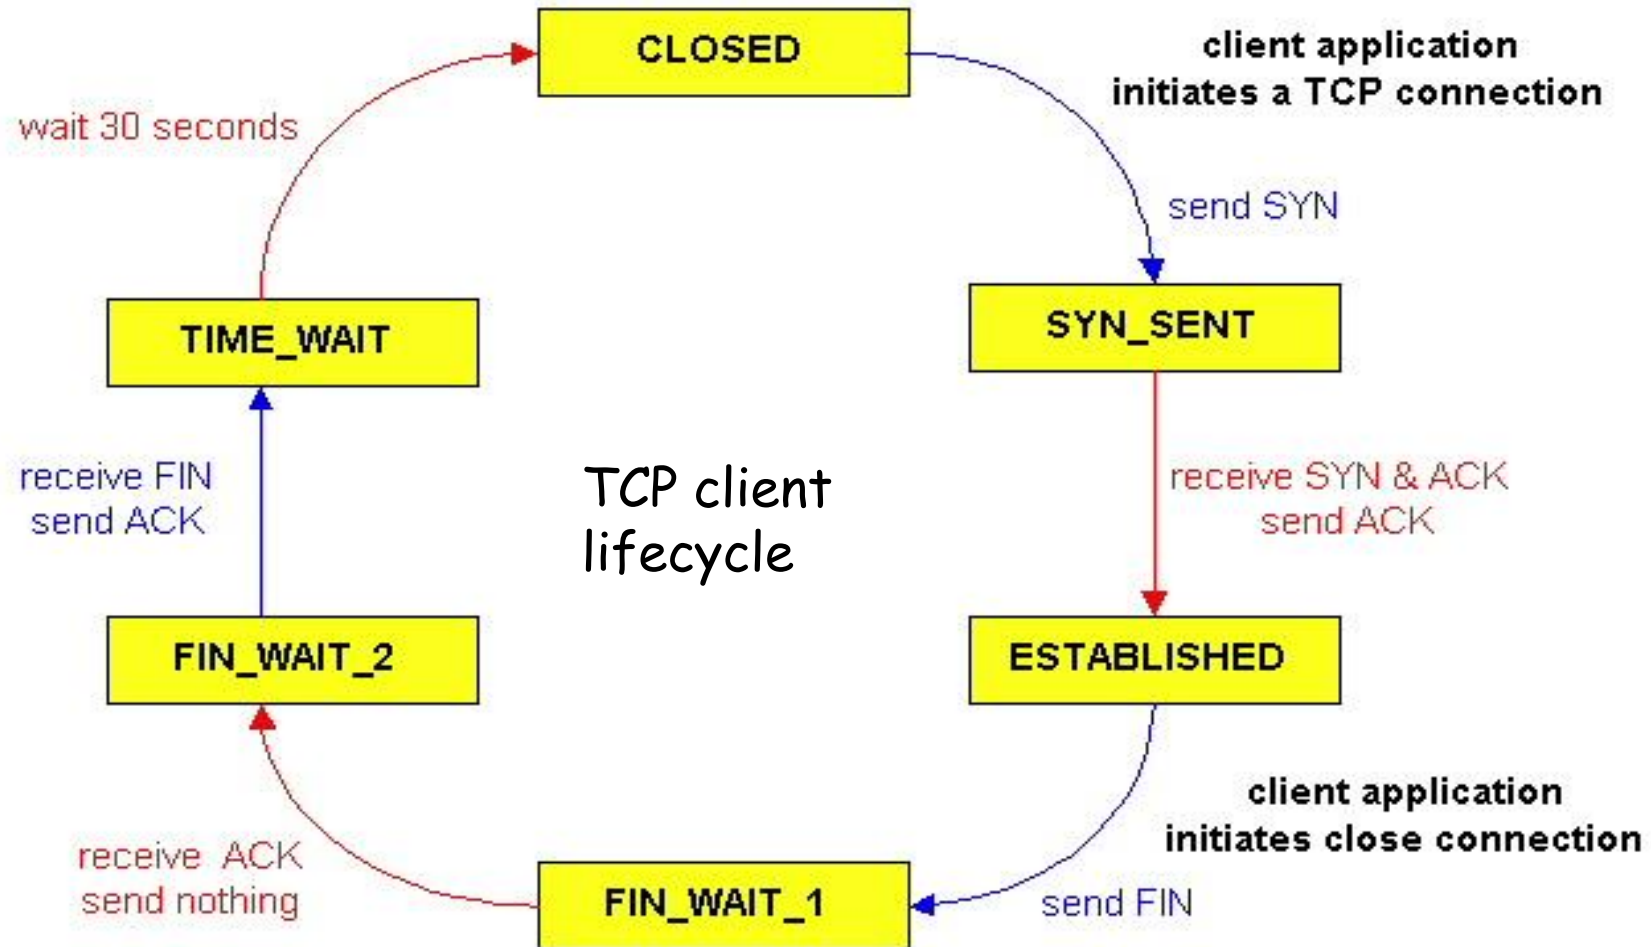
Step 6: server, receives ACK. connection fully closed

Step 7: client, timer expires, connection fully closed



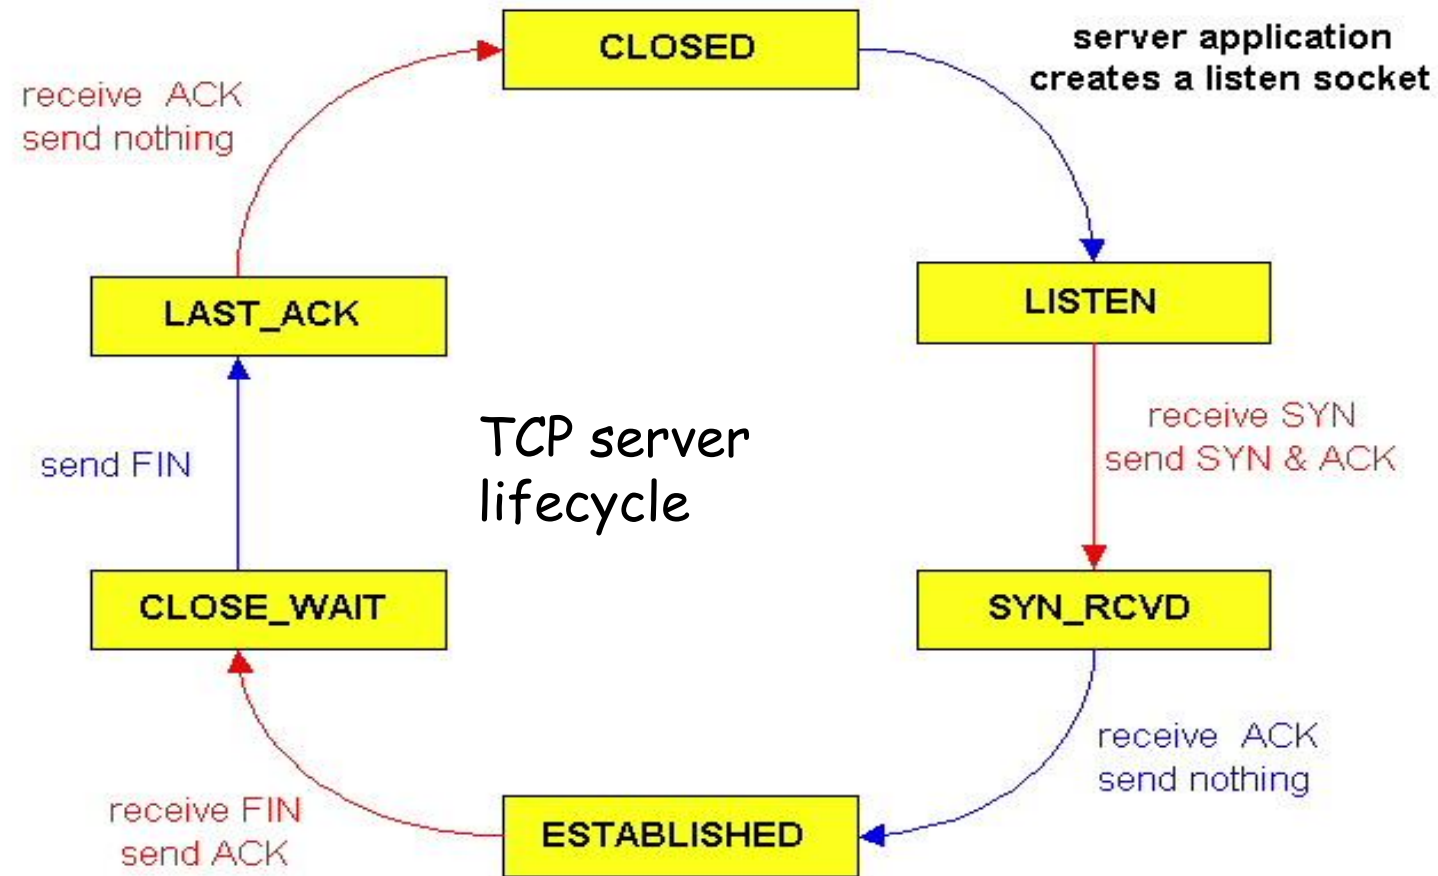
TCP Connection Management FSM

TCP client lifecycle



TCP Connection Management FSM

TCP server lifecycle

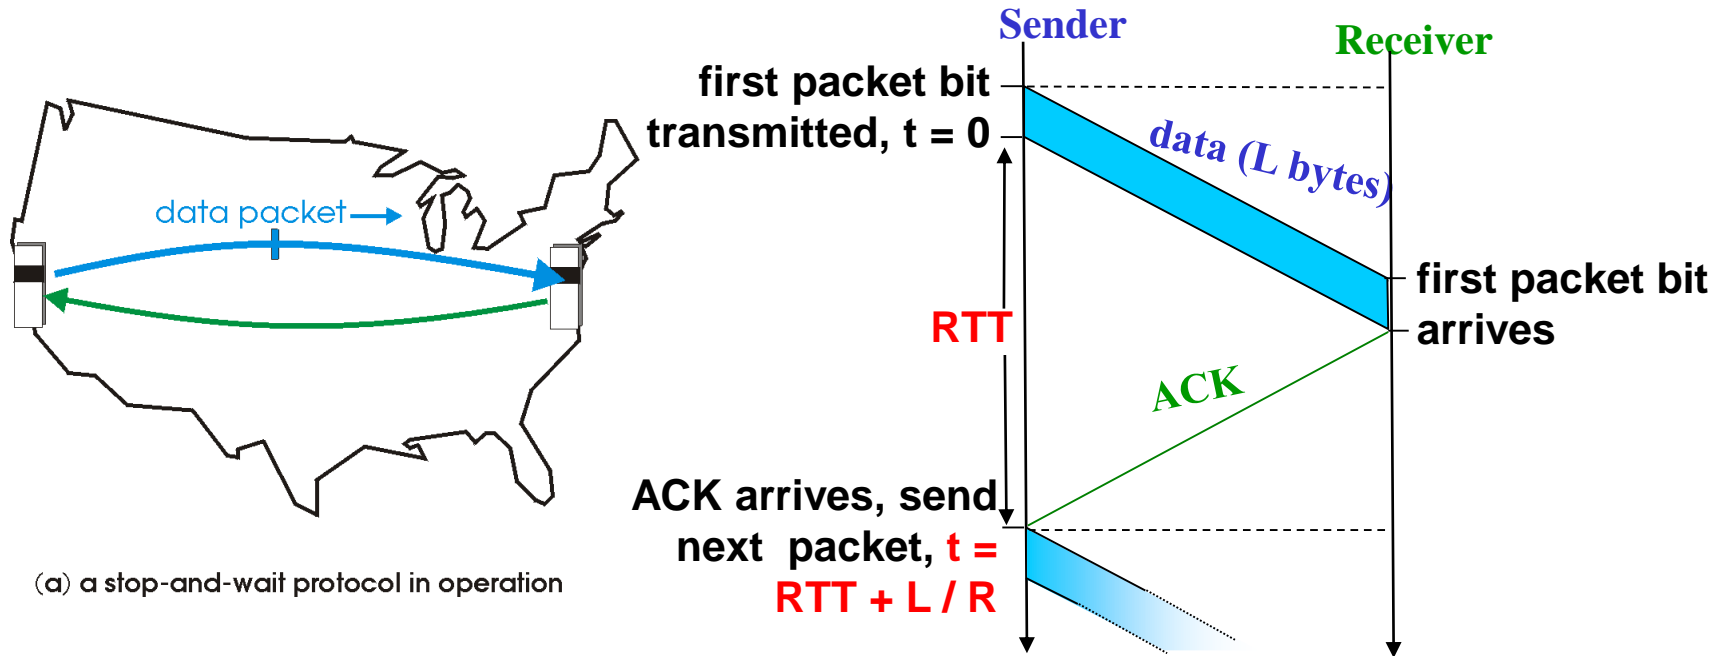


Recall:

Simple Reliable Data Transfer Protocol

- “Stop-and-Wait” Protocol
 - also called Alternating Bit Protocol
- Sender:
 - i) send data segment (n bytes) w/ $\text{seq} = x$
 - buffer data segment, set timer, retransmit if time out
 - ii) wait for ACK w/ $\text{ack} = x+n$; if received, set $x := x+n$, go to i)
 - retransmit if ACK w/ “incorrect” ack no. received
- Receiver:
 - i) expect data segment w/ $\text{seq} = x$; if received, send ACK w/ $\text{ack} = x+n$, set $x := x+n$, go to i)
 - if data segment w/ “incorrect” seq no received, discard data segment, and retransmit ACK.

Problem with Stop & Wait Protocol



- Can't keep the pipe full
 - Utilization is low
 - when bandwidth-delay product ($R \times RTT$) is large!

Stop & Wait: Performance Analysis

Example:

1 Gbps connection, 15 ms end-end prop. delay,

data segment size: 1 KB = 8Kb

$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8 \text{ kb}}{10^9 \text{ b/s}}$$
$$= 8 \times 10^{-6} \text{ s} = 0.008 \text{ ms}$$

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{L}{RTT * R + L} = \frac{.008}{30.008} = 0.00027$$

- U_{sender} : utilization, i.e., fraction of time sender busy sending
- 1KB data segment every 30 msec (round trip time)
 - > $0.027\% \times 1 \text{ Gbps} = 33\text{KB/sec}$ throughput over 1 Gbps link


Moral of story:

network protocol *limits* use of physical resources!

Pipelined Protocols

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged data segments

- range of sequence numbers must be increased
- buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- Two generic forms of pipelined protocols:
Go-Back-N and Selective Repeat

Pipelining: Increased Utilization



Increase utilization
by a factor of 3!

$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Go-Back-N: Basic Ideas

Sender:

- Packets transmitted continually (when available) without waiting for ACK, up to N outstanding, unACK'ed packets
- A logically different timer associated with each “in-flight” (i.e., unACK'ed) packet
 - *timeout(n)*: retransmit pkt n and all higher seq # pkts in window


Receiver:

- ACK packet if corrected received and in-order, pass to higher layer, NACK or ignore corrupted or out-of-order packets
- “cumulative” ACK: if multiple packets received corrected and in-order, send only one ACK with ack= next expected seq no.

Go-Back-N: Sliding Windows

Sender:

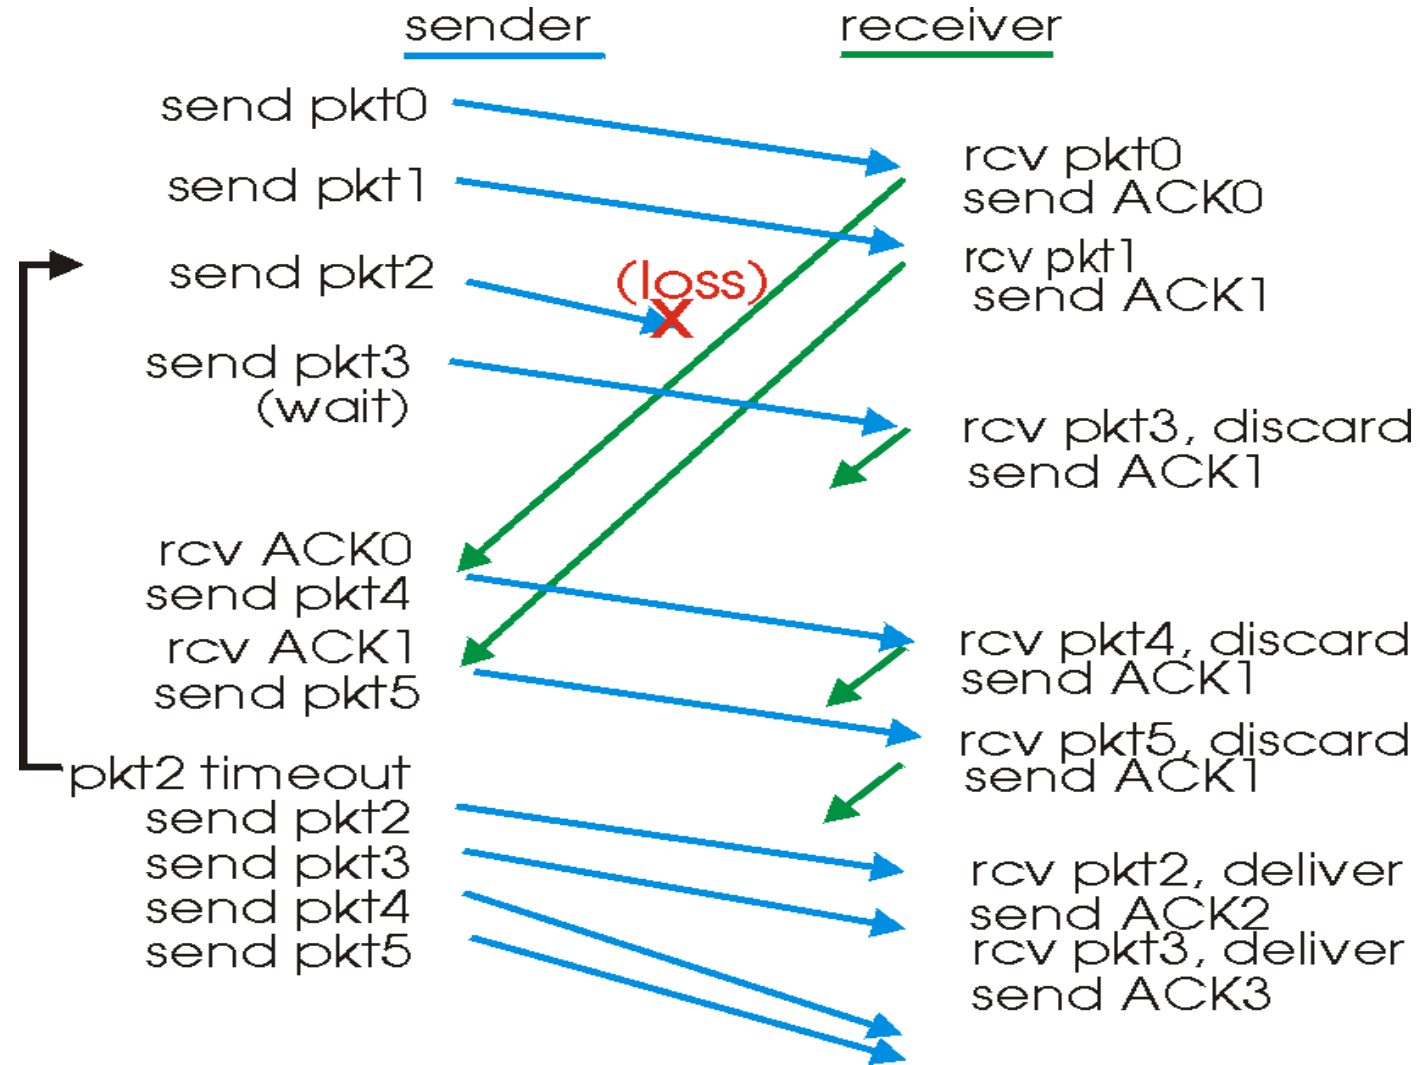
- “window” of up to N , consecutive unack’ed pkts allowed
- **send_base**: first sent but unACKed pkt, move forward when ACK’ed



Receiver:

- **rcv_base**: keep track of next expected seq no, move forward when next in-order (i.e., w/ expected seq no) pkt received

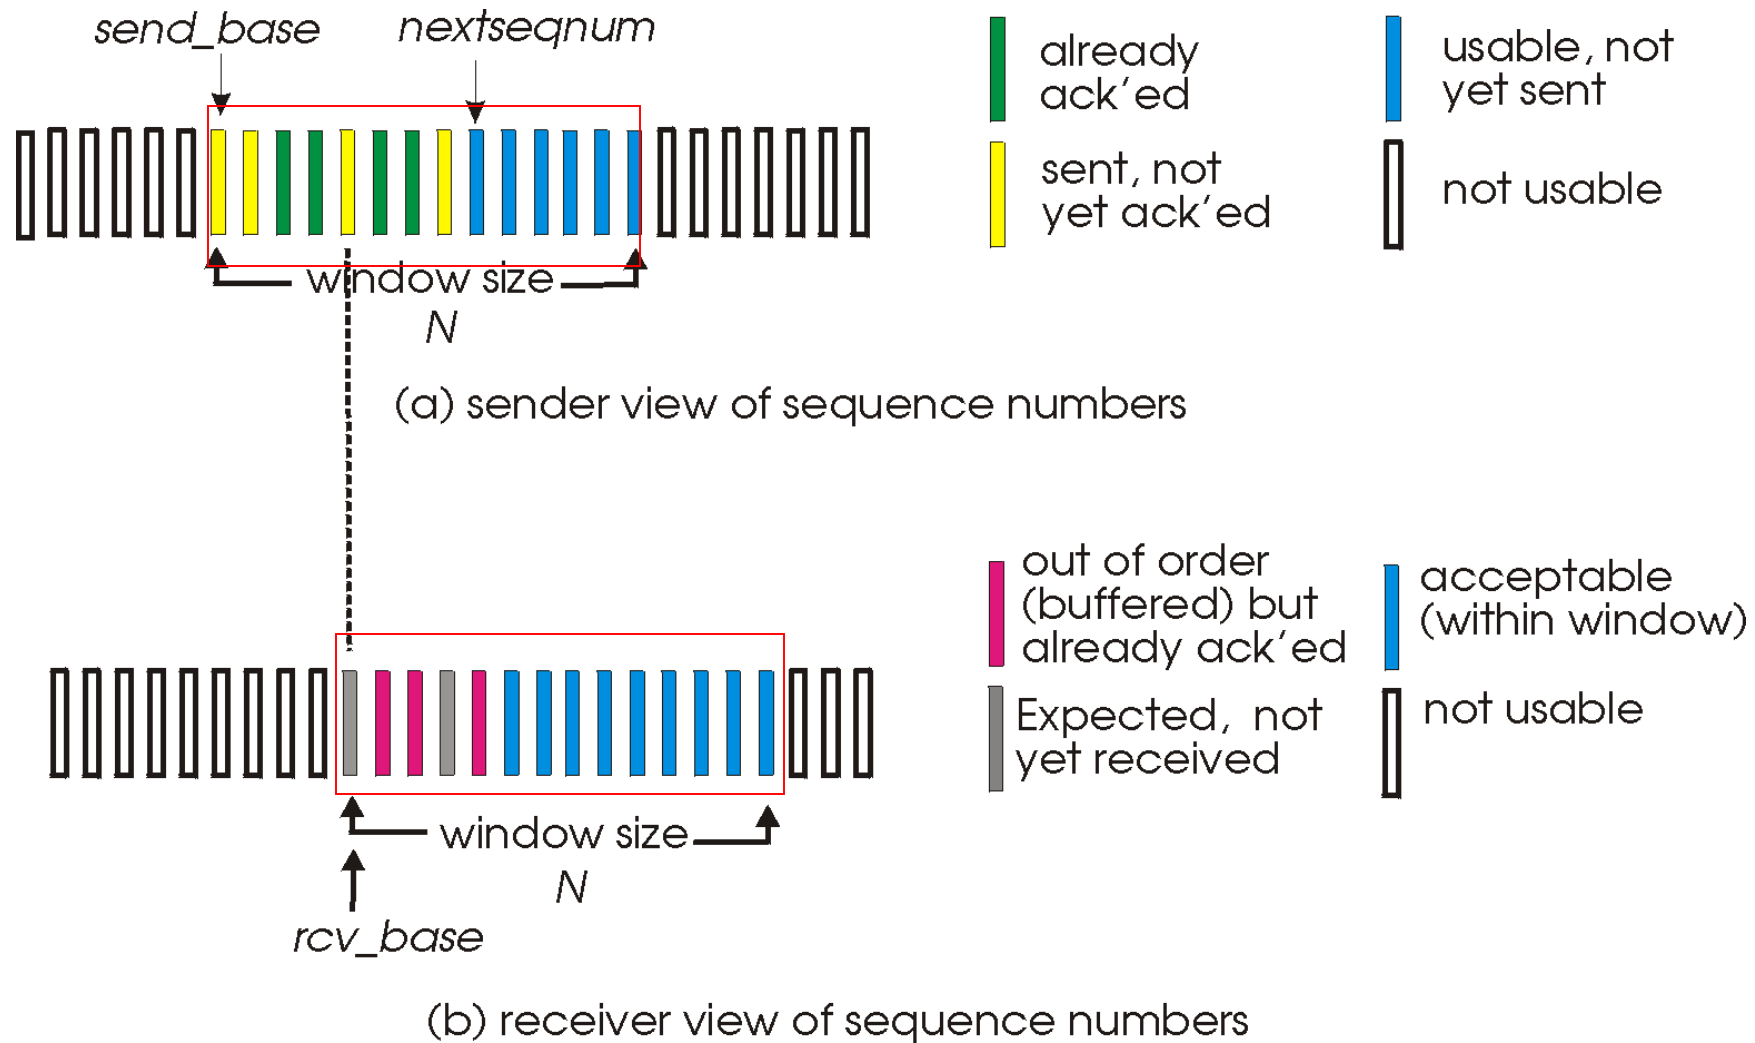
GBN in Action



Selective Repeat

- As in Go-Back-N
 - Packet sent when available up to window limit
- Unlike Go-Back-N
 - Out-of-order (but otherwise correct) is ACKed
 - Receiver: buffer out-of-order pkts, no “cumulative” ACKs
 - Sender: on timeout of packet k, retransmit just pkt k
- Comments
 - Can require more receiver buffering than Go-Back-N
 - More complicated buffer management by both sides
 - Save bandwidth
 - no need to retransmit correctly received packets

Selective Repeat: Sliding Windows



Selective Repeat: Algorithms

—sender—

data from above :

- if next available seq # in window, send pkt

timeout(n):

- resend pkt n, restart timer

ACK(n) in [sendbase, sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

—receiver—

pkt n in [rcvbase, rcvbase+N-1]

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt


pkt n in [rcvbase-N, rcvbase-1]

- ACK(n)

otherwise:

- ignore



Selective Repeat in Action



Selective Repeat: Dilemma

Example:

- seq #'s: 0, 1, 2, 3
 - window size=3
 - receiver sees no difference in two scenarios!
 - incorrectly passes duplicate data as new in (a)
- Q:** what relationship between seq # size and window size?



Seqno Space and Window Size

- How big the sliding window can be?
 - MAXSEQNO: number of available sequence numbers
 - Under Go-Back-N?
 - MAXSEQNO will not work, why?
 - What about Selective-Repeat?

TCP Round Trip Time and Timeout

Q: how to set TCP timeout value?

- longer than RTT
 - but RTT varies
- too short: premature timeout
 - unnecessary retransmissions
- too long: slow reaction to segment loss

Q: how to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
 - ignore retransmissions, why?
- **SampleRTT** will vary, want estimated RTT “smoother”
 - average several recent measurements, not just current **SampleRTT**

TCP Round Trip Time Estimation

$$\text{EstimatedRTT} = (1-a) * \text{EstimatedRTT} + a * \text{SampleRTT}$$

- Exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $a = 0.125$

Setting the timeout interval

- EstimatedRTT plus “safety margin”
 - large variation in EstimatedRTT -> larger safety margin
- “safety margin”: accommodate variations in estimatedRTT


$$\text{DevRTT} = (1-b) * \text{DevRTT} + b * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $b = 0.25$)

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

Example RTT Estimation:

RTT: gaia.cs.umass.edu to fantasia.eurecom.fr



TCP Reliable Data Transfer

- TCP creates reliable data transfer service on top of IP's unreliable service
- Pipelined segments
- Cumulative ACKs
- TCP uses single retransmission timer
 - double TimeoutInterval on timer expiration
- Retransmissions are triggered by:
 - timeout events
 - duplicate acks
- Initially consider simplified TCP sender:
 - ignore duplicate acks
 - ignore flow control, congestion control

TCP Sender Events:

data rcvd from app:

- Create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running (think of timer as for oldest unacked segment)
- expiration interval:
TimeOutInterval

timeout:

- retransmit segment that caused timeout
- restart timer

ACK received:

- If acknowledges previously unACKed segments, then
 - update what is known to be ACKed
 - start timer if there are outstanding segments

TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver	TCP Receiver Action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected	Immediately send duplicate ACK, indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap